

```
#Name:Yash Bhupesh Aware
```

```
#Roll no. :02
```

```
#Sub :AIML
```

```
#section:3C
```

```
#Date:23/01/2025
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
import os
```

```
os.getcwd()
```

```
'C:\\Users\\PC-6'
```

```
os.chdir("C:\\Users\\PC-6\\Desktop")
```

```
df=pd.read_csv("Housing.csv")
```

```
df.head()
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom
basement \							
0	13300000	7420	4	2	3	yes	no
no							
1	12250000	8960	4	4	4	yes	no
no							
2	12250000	9960	3	2	2	yes	no
yes							
3	12215000	7500	4	2	2	yes	no
yes							
4	11410000	7420	4	1	2	yes	yes
yes							

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

```
df.isna().sum()
```

price	0
area	0
bedrooms	0
bathrooms	0
stories	0

```
mainroad      0
guestroom     0
basement      0
hotwaterheating 0
airconditioning 0
parking       0
prefarea      0
furnishingstatus 0
dtype: int64
```

```
df['furnishingstatus'].nunique()
```

```
3
```

```
columns_to_transform = ['mainroad', 'guestroom',
                        'basement', 'hotwaterheating', 'airconditioning', 'prefarea']
df[columns_to_transform] = df[columns_to_transform].replace({'yes': 1,
                                                            'no': 0})
```

```
df['furnishingstatus'] =
df['furnishingstatus'].replace({'unfurnished': 0, 'semi-furnished': 1,
                                'furnished': 2})
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
sc = ['price', 'area']
df[sc] = scaler.fit_transform(df[sc])
```

```
df
```

	price	area	bedrooms	bathrooms	stories	mainroad
guestroom \						
0	4.566365	1.046726	4	2	3	1
0						
1	4.004484	1.757010	4	4	4	1
0						
2	4.004484	2.218232	3	2	2	1
0						
3	3.985755	1.083624	4	2	2	1
0						
4	3.554979	1.046726	4	1	2	1
1						
..
...						
540	-1.576868	-0.991879	2	1	1	1
0						
541	-1.605149	-1.268613	3	1	1	0
0						
542	-1.614327	-0.705921	2	1	1	1
0						

```

543 -1.614327 -1.033389      3      1      1      0
0
544 -1.614327 -0.599839      3      1      2      1
0

```

	basement	hotwaterheating	airconditioning	parking	prefarea	\
0	0	0	1	2	1	
1	0	0	1	3	0	
2	1	0	0	2	1	
3	1	0	1	3	1	
4	1	0	1	2	0	
..	
540	1	0	0	2	0	
541	0	0	0	0	0	
542	0	0	0	0	0	
543	0	0	0	0	0	
544	0	0	0	0	0	

```

furnishingstatus
0      2
1      2
2      1
3      2
4      2
..      ...
540     0
541     1
542     0
543     2
544     0

```

[545 rows x 13 columns]

df.dtypes


```

price      float64
area       float64
bedrooms   int64
bathrooms  int64
stories    int64
mainroad   int64
guestroom  int64
basement   int64
hotwaterheating int64
airconditioning int64
parking     int64
prefarea    int64
furnishingstatus int64
dtype: object

```

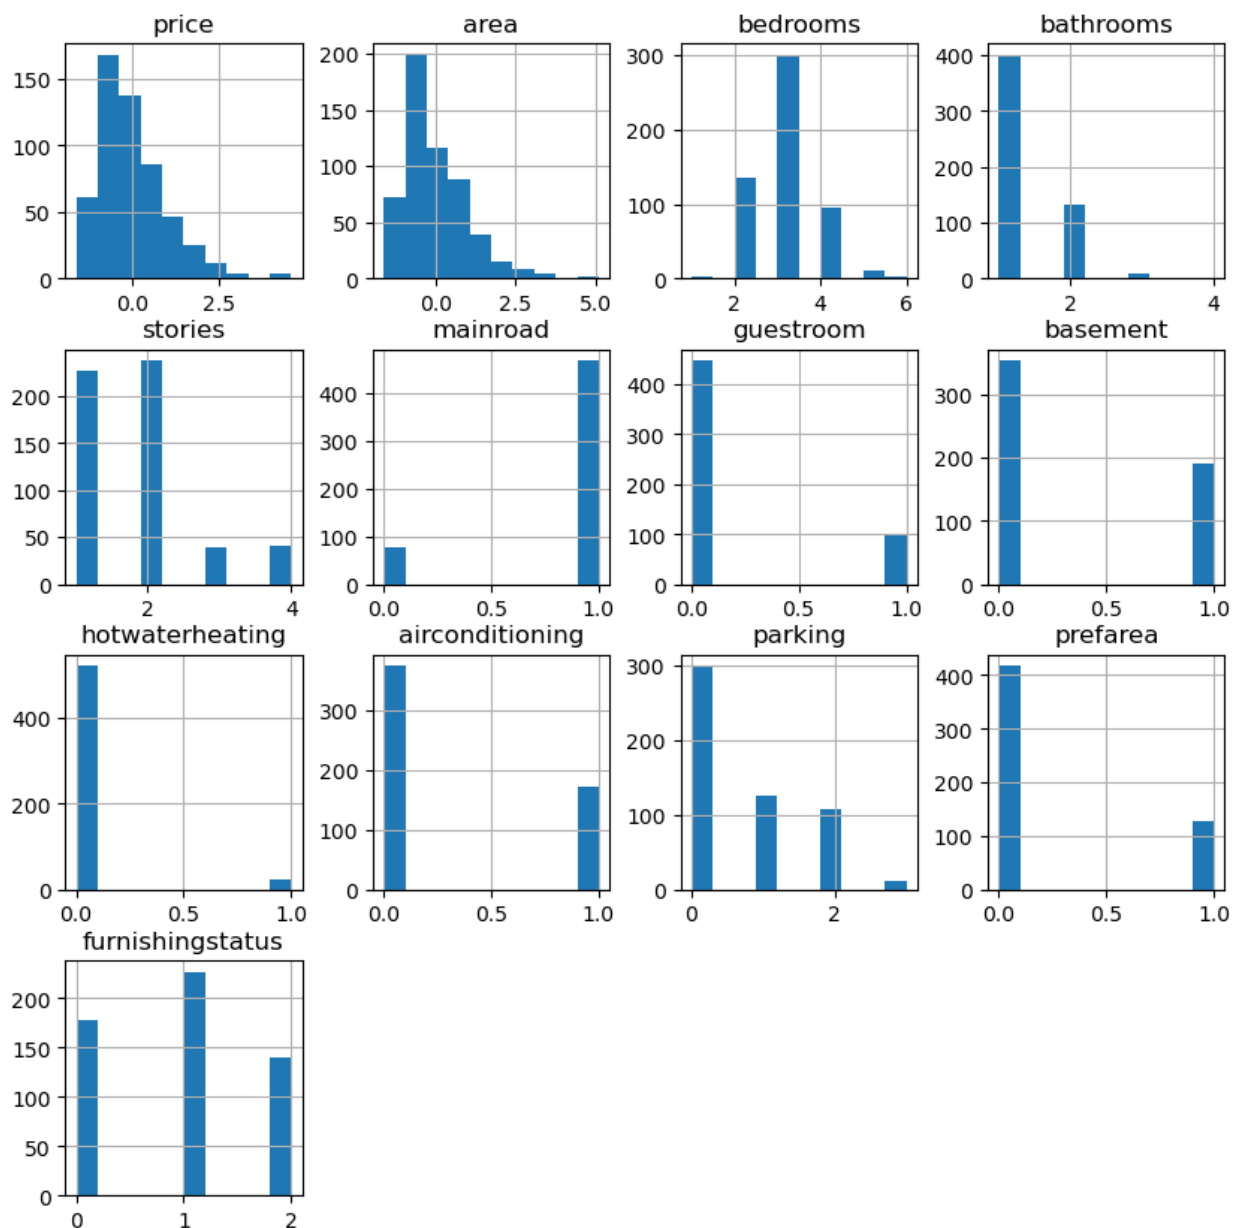
```
corr_matrix = df.corr()

plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm")
plt.show()
```



```
df.hist(figsize=(10, 10), bins=10)
plt.suptitle("Histograms for All Columns", fontsize=16)
plt.show()
```

Histograms for All Columns




```
X = df.drop('price', axis=1)
y = df['price']

X_features = ['area', 'bedrooms', 'bathrooms']

fig, ax = plt.subplots(1, 3, figsize=(12, 3), sharey=True)
for i in range(3): # Assuming there are 4 features
    ax[i].scatter(X.iloc[:, i], y)
    ax[i].set_xlabel(X_features[i])
```

```
ax[0].set_ylabel("Price z-score Normalized")
plt.show()
```



```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

from sklearn.linear_model import LinearRegression

lr_model = LinearRegression()
lr_model.fit(X_train, y_train)

y_pred = lr_model.predict(X_test)

from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV

dt_model = DecisionTreeRegressor(random_state=42)

param_grid = {
    'max_depth': [5, 10, 15, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 5, 10],
    'max_features': ['auto', 'sqrt', 'log2'],
    'criterion': ['squared_error', 'friedman_mse', 'poisson',
'absolute_error']
}

grid_search = GridSearchCV(estimator=dt_model, param_grid=param_grid,
cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

print("Best Parameters:", grid_search.best_params_)

best_dt_model = grid_search.best_estimator_

Best Parameters: {'criterion': 'absolute_error', 'max_depth': 15,
'max_features': 'sqrt', 'min_samples_leaf': 5, 'min_samples_split': 2}
```

```
lr_y_pred = lr_model.predict(X_test)
dt_y_pred = best_dt_model.predict(X_test)

from sklearn.metrics import r2_score
lr_acc = r2_score(y_test, lr_y_pred)
dt_acc = r2_score(y_test, dt_y_pred)


print(f"Linear Regression Accuracy: {lr_acc}")
print(f"Decision Tree Accuracy: {dt_acc}")

Linear Regression Accuracy: 0.6494754192267798
Decision Tree Accuracy: 0.4928485257199178

plt.figure(figsize=(8, 6))
plt.scatter(y_test, lr_y_pred, color='blue', alpha=0.7,
            label='Predictions')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
         color='red', linestyle='--', label='Perfect Fit')

plt.xlabel('Target Values')
plt.ylabel('Predicted Values')
plt.title('Linear Regression Predicted vs Target Values')
plt.legend()
plt.grid(True)


plt.show()
```



```
plt.figure(figsize=(8, 6))
plt.scatter(y_test, dt_y_pred, color='blue', alpha=0.7,
            label='Predictions')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
         color='red', linestyle='--', label='Perfect Fit')

plt.xlabel('Target Values')
plt.ylabel('Predicted Values')
plt.title('Decision Tree Predicted vs Target Values')
plt.legend()
plt.grid(True)

plt.show()
```

```

from sklearn.linear_model import SGDRegressor

sgdr = SGDRegressor(max_iter=1000)
sgdr.fit(X_train, y_train)
print(sgdr)
print(f"number of iterations completed: {sgdr.n_iter_}, number of
weight updates: {sgdr.t_}")

SGDRegressor()
number of iterations completed: 34, number of weight updates: 14825.0

b_norm = sgdr.intercept_
w_norm = sgdr.coef_
print(f"model parameters:                w: {w_norm}, b:{b_norm}")

model parameters:                w: [ 0.33710478 -0.14391223
0.44655225  0.22453196 -0.14067121  0.13448173
0.20822998  0.18997849  0.43576092  0.12443931  0.32087466
0.09175696], b: [-0.98702889]

```

```

y_pred_sgd = sgdr.predict(X_train)
y_pred = np.dot(X_train, w_norm) + b_norm
print(f"prediction using np.dot() and sgdr.predict match: {(y_pred ==
y_pred_sgd).all()}")

print(f"Prediction on training set:\n{y_pred[:4]}" )
print(f"Target values \n{y_train[:4]}")

prediction using np.dot() and sgdr.predict match: True
Prediction on training set:
[ 1.1075839  0.98591619  0.12444803 -0.52972045]
Target values
46      1.476019
93      0.820491
335     -0.453106
412     -0.715317
Name: price, dtype: float64

fig,ax=plt.subplots(1,3,figsize=(12,3),sharey=True)
for i in range(len(ax)):
    ax[i].scatter(X_train.iloc[:,i],y_train, label = 'target')
    ax[i].set_xlabel(X_features[i])
    ax[i].scatter(X_train.iloc[:,i],y_pred,color="orange", label =
'predict')
ax[0].set_ylabel("Price"); ax[0].legend();
fig.suptitle("target versus prediction using z-score normalized
model")
plt.show()

```

